

1. Remember: Core STM Concepts and Concurrency Issues

Define **TVar** and Basic **STM** Operations

```
...  
  
import Control.Concurrent.MVar  
import Control.Concurrent  
  
-- Define a transactional variable (TVar)  
data TVar a = TVar (MVar a)  
  
-- Function to create a new TVar  
newTVar :: a -> IO (TVar a)  
newTVar val = TVar <$> newMVar val  
  
-- Function to read from a TVar  
readTVar :: TVar a -> IO a  
readTVar (TVar mvar) = do  
    val <- readMVar mvar  
    threadDelay 5000000 -- Add a delay of 5 second  
    return val  
  
-- Function to write to a TVar  
writeTVar :: TVar a -> a -> IO ()  
writeTVar (TVar mvar) val = modifyMVar_ mvar (\_ -> return val)  
...
```

Identifying concurrency issues: This simple version of **TVar** does not yet solve concurrency problems like race conditions or deadlocks. It uses **MVar** for thread safety but does not handle transaction rollbacks or retries.

Atomicity, Consistency, Isolation: The next steps will aim to build the transaction mechanism to achieve atomicity and isolation by running multiple operations as a single unit (the core of STM).

As an experiment, we can proceed as follows:

1. Define thread variable **tv** with a default value.
2. Read and double check the value
3. Have a **readTVar** func running in background
4. Attempt to **writeTVar** during the background execution
5. Read values to double check.

As a result, everything plays out as expected, the function reads the value before the value gets overwritten.

```

ghci> :l main1.hs
[1 of 2] Compiling Main                ( main1.hs, interpreted )
Ok, one module loaded.
ghci> tv <- newTVar 0
ghci> readTVar tv >=> print
0
ghci> forkIO (readTVar tv >=> print)
ThreadId 868
ghci> writeTVar tv 100
ghci> 0

ghci> readTVar tv >=> print
100
ghci> █

```

2. Understand: STM for Preventing Concurrency Issues

Next, I implement a very basic version of STM to group operations on **TVars** within an atomic block. This will help prevent inconsistent states when multiple threads access shared resources concurrently.

Implement a Simple STM Monad

...

```

import Control.Exception (try, SomeException)
import Control.Concurrent.MVar
import Control.Concurrent

-- Define a transactional variable (TVar)
data TVar a = TVar (MVar a)

-- Function to create a new TVar
newTVar :: a -> IO (TVar a)
newTVar val = TVar <$> newMVar val

-- Function to read from a TVar
readTVar :: TVar a -> IO a
readTVar (TVar mvar) = do
    val <- readMVar mvar
    threadDelay 15000000 -- Add a delay of 5 second

```

```

        return val

-- Function to write to a TVar
writeTVar :: TVar a -> a -> IO ()
writeTVar (TVar mvar) val = modifyMVar_ mvar (\_ -> return val)

-- Stage 2 update
-- Define a simple STM monad
newtype STM a = STM { runSTM :: IO a }

-- The atomically function to execute STM actions
atomically :: STM a -> IO a
atomically (STM action) = do
    -- Wrap action in exception handling to simulate a retry mechanism
    result <- try action
    case result of
        Right val -> return val
        Left (_ :: SomeException) -> atomically (STM action)  -- Retry on
failure

-- Functions to lift TVar operations into STM
readTVarSTM :: TVar a -> STM a
readTVarSTM tvar = STM $ readTVar tvar

writeTVarSTM :: TVar a -> a -> STM ()
writeTVarSTM tvar val = STM $ writeTVar tvar val
...

```

Here, I introduced the **STM** monad, which wraps IO operations on **TVars**. The **atomically** function executes **STM** transactions, and in case of an exception (such as a conflict in a real STM system), it retries the operation.

- **Concurrency issues and STM's retry mechanism:** We simulate the retry mechanism here by catching exceptions, but a real implementation would involve detecting conflicts between transactions.

We can repeat the same experiment as mentioned in the previous stage which yields the same expected result.

```

ghci> :l main2.hs
[1 of 2] Compiling Main                ( main2.hs, interpreted )
Ok, one module loaded.
ghci> tv <- newTVar 0
ghci> tv <- newTVar 1
ghci> atomically (readTVarSTM tvar) >>= print

<interactive>:53:25: error: Variable not in scope: tvar :: TVar a0
ghci> atomically (readTVarSTM tv) >>= print
1
ghci> forkIO (atomically (readTVarSTM tv) >>= print)
ThreadId 1566
ghci> atomically (writeTVarSTM tv 100)
ghci> 1

ghci>
ghci>
ghci> atomically (readTVarSTM tv) >>= print
100

```

3. Apply: Implement a Job Scheduler Using the Custom STM

Next, I built a simple job queue using the [STM](#) and [TVar](#) implementations which will process jobs concurrently while using STM to ensure safe access to shared data.

```

...
type JobQueue = TVar [String]

-- Submit a job to the queue
submitJobSTM :: JobQueue -> String -> STM ()
submitJobSTM queue job = do
    jobs <- readTVarSTM queue
    writeTVarSTM queue (jobs ++ [job])

-- Process a job from the queue
processJobSTM :: JobQueue -> STM (Maybe String)
processJobSTM queue = do
    jobs <- readTVarSTM queue
    case jobs of
        [] -> return Nothing

```

```

    (job:rest) -> do
        writeTVarSTM queue rest
        return (Just job)

-- Concurrently process jobs using custom STM
processJobsConcurrentlySTM :: JobQueue -> IO ()
processJobsConcurrentlySTM queue = atomically $ do
    mJob <- processJobSTM queue
    case mJob of
        Just job -> STM $ putStrLn ("Processing job: " ++ job)
        Nothing   -> return ()
    ...

```

This job scheduler submits and processes jobs using the custom **TVar** and **STM**. It ensures safe concurrent access by grouping **readTVar** and **writeTVar** operations into atomic transactions

To test the job queue, I also added a test case to analyze how it works.

```

...

-- Test case to test concurrency in the job queue
testConcurrentJobProcessing :: IO ()
testConcurrentJobProcessing = do
    -- Step 1: Create an empty job queue
    queue <- newTVar []

    -- Step 2: Spawn multiple threads to concurrently submit jobs
    let numJobs = 10 -- Total number of jobs to submit
    putStrLn "Submitting jobs concurrently..."
    replicateM_ 10 $ forkIO $ replicateM_ (numJobs `div` 10) $ do
        atomically $ submitJobSTM queue "TestJob"

    -- Step 3: Spawn multiple threads to process jobs concurrently
    putStrLn "Processing jobs concurrently..."
    replicateM_ 5 $ forkIO $ replicateM_ (numJobs `div` 5) $
        processJobsConcurrentlySTM queue

    -- Step 4: Allow time for job submission and processing
    threadDelay 2000000 -- Wait for 2 seconds to allow all jobs to be
        processed

    -- Step 5: Check if all jobs were processed

```

```

jobsRemaining <- readTVar queue
putStrLn $ "Remaining jobs: " ++ show jobsRemaining

-- Since we've submitted `numJobs` jobs, the queue should be empty
after all are processed
when (null jobsRemaining) $
    putStrLn "All jobs processed. PASS"
when (not (null jobsRemaining)) $
    putStrLn "Some jobs were not processed. FAIL"
...

```

And here is the result of it:

```

ghci> testConcurrentJobProcessing
Submitting jobs concurrently...
Submitting jobs concurrently...
Processing jobs concurrently...
Processing jobs concurrently...
PrPPoPrPrCrroeoecocsccecesesesissnsisiginin ngngjg g o j jbjjo:obob b:b:T: : e T TsTeTetesesJststotJtJ
bJoJo
obobb
b

Remaining jobs: []
All jobs processed. PASS
ghci>

```

4. Analyze: Implementing and Benchmarking Concurrency

To analyze the performance of my custom implementation, I implemented two versions of the concurrency code, one with custom **STM** and one with **MVar**.

The main difference between is how the JobQueue is instantiated

CustomSTM	MVar
<code>type JobQueue = TVar [String]</code>	<code>type JobQueue = MVar [String]</code>

After a couple of experiments, there is no significant difference between the two implementations. As expected, **CustomSTM** performs worse in time (by 0.2%).

Here are the results of the test (**CustomSTM** above, **MVar** below):

